

# **Database Management System**

**UNIT II: Chapter 6: Concurrency Control Techniques**

# Lock and unlock operations for binary locks

**lock\_item( $X$ ):**

**B:**   if  $\text{LOCK}(X) = 0$                                    (\* item is unlocked \*)  
          then  $\text{LOCK}(X) \leftarrow 1$            (\* lock the item \*)  
      else  
          **begin**  
          wait (until  $\text{LOCK}(X) = 0$   
              and the lock manager wakes up the transaction);  
          go to **B**  
          **end;**

**unlock\_item( $X$ ):**

$\text{LOCK}(X) \leftarrow 0$ ;                                   (\* unlock the item \*)  
      if any transactions are waiting  
          then wakeup one of the waiting transactions;



## **In binary locks every transaction must obey the following rules:**

1. A transaction T must issue the operation `lock_item(X)` before any `read_item(X)` or `write_item(X)` operations are performed in T.
2. A transaction T must issue the operation `unlock_item(X)` after all `read_item(X)` and `write_item(X)` operations are completed in T.
3. A transaction T will not issue a `lock_item(X)` operation if it already holds the lock on item X.
4. A transaction T will not issue an `unlock_item(X)` operation *unless it already* holds the lock on item X.



# Locking and unlocking operations for two mode locks (Read/Write Locks)

**read\_lock( $X$ ):**

**B:** if  $\text{LOCK}(X) = \text{"unlocked"}$

    then **begin**  $\text{LOCK}(X) \leftarrow \text{"read-locked"}$ ;

$\text{no\_of\_reads}(X) \leftarrow 1$

**end**

else if  $\text{LOCK}(X) = \text{"read-locked"}$

    then  $\text{no\_of\_reads}(X) \leftarrow \text{no\_of\_reads}(X) + 1$

else **begin**

    wait (until  $\text{LOCK}(X) = \text{"unlocked"}$

        and the lock manager wakes up the transaction);

    go to **B**

**end;**



---

**write\_lock( $X$ ):**

**B:** if LOCK( $X$ ) = "unlocked"

    then LOCK( $X$ )  $\leftarrow$  "write-locked"

    else **begin**

        wait (until LOCK( $X$ ) = "unlocked"

            and the lock manager wakes up the transaction);

        go to **B**

**end;**



---

**unlock (X):**

    if  $\text{LOCK}(X) = \text{"write-locked"}$

        then **begin**  $\text{LOCK}(X) \leftarrow \text{"unlocked"}$ ;

            wakeup one of the waiting transactions, if any

**end**

    else if  $\text{LOCK}(X) = \text{"read-locked"}$

        then **begin**

$\text{no\_of\_reads}(X) \leftarrow \text{no\_of\_reads}(X) - 1$ ;

            if  $\text{no\_of\_reads}(X) = 0$

                then **begin**  $\text{LOCK}(X) = \text{"unlocked"}$ ;

                    wakeup one of the waiting transactions, if any

**end**

**end**;



# Shared/Exclusive locking scheme

**When we use the shared/exclusive locking scheme, the system must enforce the following rules:**

1. A transaction T must issue the operation `read_lock(X)` or `write_lock(X)` before any `read_item(X)` operation is performed in T.
2. A transaction T must issue the operation `write_lock(X)` before any `write_item(X)` operation is performed in T.
3. A transaction T must issue the operation `unlock(X)` after all `read_item(X)` and `write_item(X)` operations are completed in T.



## Shared/Exclusive locking scheme

4. A transaction T will not issue a `read_lock(X)` operation if it already holds a read (shared) lock or a write (exclusive) lock on item X. This rule may be **relaxed**.
5. A transaction T will not issue a `write_lock(X)` operation if it already holds a read (shared) lock or write (exclusive) lock on item X. This rule may be **relaxed**.
6. A transaction T will not issue an `unlock(X)` operation unless it already holds a read (shared) lock or a write (exclusive) lock on item X.





# Does not guarantee serializability

(a)

$T_1$	$T_2$
<pre>read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X);</pre>	<pre>read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y);</pre>

(b)

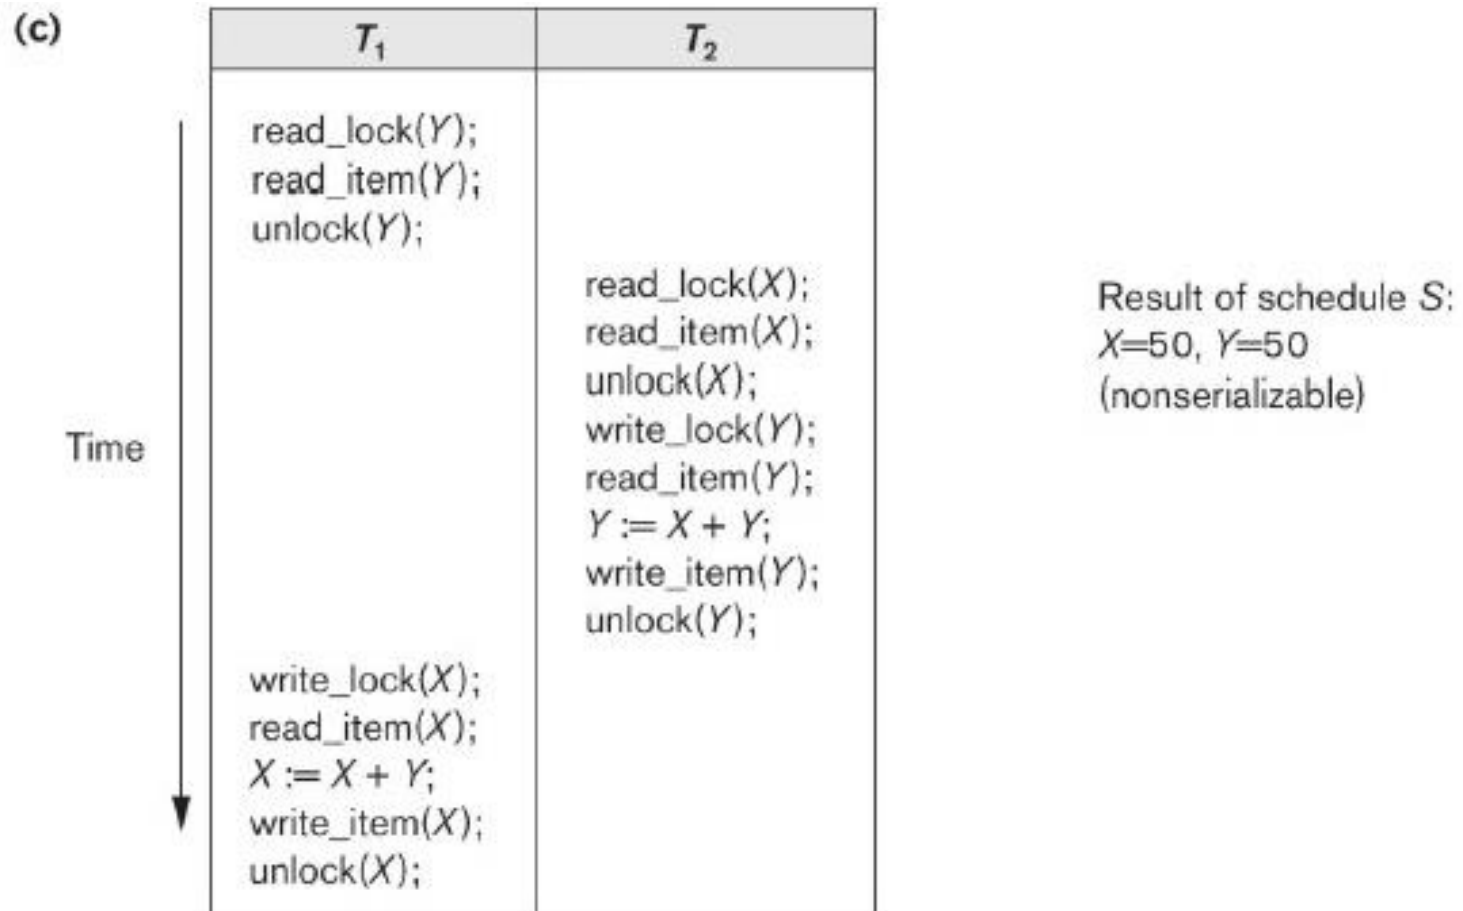
Initial values:  $X=20, Y=30$

Result serial schedule  $T_1$   
followed by  $T_2$ :  $X=50, Y=80$

Result of serial schedule  $T_2$   
followed by  $T_1$ :  $X=70, Y=50$



# A nonserializable schedule using locks



## 2PL [schedule is guaranteed to be serializable]

$T_1'$	$T_2'$
<pre>read_lock(Y); read_item(Y); write_lock(X); unlock(Y) read_item(X); X := X + Y; write_item(X); unlock(X);</pre>	<pre>read_lock(X); read_item(X); write_lock(Y); unlock(X) read_item(Y); Y := X + Y; write_item(Y); unlock(Y);</pre>



# Conversion of Locks: Relax 4 and 5

- Ⓒ upgrade the lock

- Ⓒ downgrade the lock

- Ⓒ If lock conversion is allowed in 2PL, then

  - upgrading of locks (from read-locked to write-locked) must be done during the **expanding phase**.

  - downgrading of locks (from write-locked to read-locked) must be done in the **shrinking phase**.



# 2PL and Serializability

Ⓒ Two-phase locking protocol guarantees serializability.

Ⓒ The use of locks can cause two additional problems:

1. Deadlock and
2. Starvation



# Variations of 2PL

---

**Basic,  
Conservative,  
Strict,  
Rigorous**



# Conservative 2PL

A variation known as **conservative 2PL** (or **static 2PL**) requires a transaction to **lock** all the items it accesses before the transaction begins execution, by predeclaring its read-set and write-set.

If any of the predeclared items needed cannot be locked, the transaction does not lock any item; instead, it **waits** until all the items are available for locking.

Conservative 2PL is a **deadlock-free** protocol.



# Strict 2PL

---

A transaction T does not release any of its **exclusive (write) locks** until after it commits or aborts.

Hence, no other transaction can read or write an item that is written by T unless T has committed, leading to a **strict schedule** for recoverability.

Strict 2PL is **not deadlock-free**.





# Rigorous 2PL

In **Rigorous 2PL** a transaction T does not release any of its **locks (exclusive or shared)** until after it commits or aborts.

Rigorous 2PL. also guarantees **strict schedules**.



# Conservative v.s. Rigorous 2PL

**Conservative 2PL** must lock all its items before it starts so once the transaction starts it is in its **shrinking phase**.

**Rigorous 2PL** does not unlock any of its items until after it terminates (by committing or aborting) so the transaction is in its **expanding phase** until it ends.



# 2PL Example

---

Consider the following two transactions:

<b>T1:</b> read(A); read(B); if then B := B + 1; write(B).	<b>T2:</b> read(B); read(A); if then A := A + 1; write(A).
---	---

Add lock and unlock instructions to transactions T1 and T2, so that they observe the two-phase locking protocol. Can the execution of these transactions result in a deadlock?



## 2PL Example

Consider the following two transactions:

T1 = w1(C) r1(A) w1(A) r1(B) w1(B)

T2 = r2(B) w2(B) r2(A) w2(A)

Assume that the scheduler uses exclusive locks only. For each of the following instances involving T1 and T2, annotated with lock and unlock actions.



# 2PL Example cont...

## **Instance A**

T1= L1(C) w1(C) L1(A) r1(A) w1(A) U1(A) L1(B) r1(B) w1(B)  
U1(C) U1(B)

T2= L2(B) r2(B) w2(B) U2(B) L2(A) r2(A) w2(A) U2(A)

## **Instance B**

T1= L1(C) w1(C) L1(A) r1(A) w1(A) L1(B) r1(B) w1(B)  
COMMIT U1(A) U1(C) U1(B)

T2= L2(B) r2(B) w2(B) L2(A) r2(A) w2(A) COMMIT U2(A)  
U2(B)

**Is it Serializable, 2PL, Strict 2PL?**

---



# 2PL Example cont...

## **Instance C**

T1= L1(C) L1(A) w1(C) r1(A) w1(A) L1(B) r1(B) w1(B)  
COMMIT U1(A) U1(C) U1(B)

T2= L2(B) r2(B) w2(B) L2(A) r2(A) w2(A) COMMIT U2(A)  
U2(B)

## **Instance D**

T1= L1(B) L1(C) w1(C) L1(A) r1(A) w1(A) r1(B) w1(B)  
COMMIT U1(A) U1(C) U1(B)

T2= L2(B) r2(B) w2(B) L2(A) r2(A) w2(A) COMMIT U2(A)  
U2(B)

**Is it Serializable, 2PL, Strict 2PL?**

---



# Timestamping

Timestamp: a unique identifier created by DBMS that indicates relative starting time of a transaction.

Can be generated by:

- using system clock at time transaction started, or
- incrementing a logical counter every time a new transaction starts.



# Timestamp based concurrency control

Does not use any locks for the serializability of schedules.

Hence, **no deadlocks**

Each transaction will be assigned a timestamp value, **TS(T<sub>i</sub>)** and based upon this the transactions are executed

The timestamp ordering protocol ensures that any **conflicting read and write operations** are executed in timestamp order.

This order decides proper **serializability** of schedules





# Notations

---

**W-TS(X):** is the largest timestamp of any transaction that executed W(X) successfully

**R-TS(X):** is the largest timestamp of any transaction that executed R(X) successfully

**TS(Ti):** Timestamp value of Transaction Ti



# Basic Time stamp protocol/alg

```
if  $T_i$  requests read( $X$ ) { if
     $TS(T_i) < W-TS(X)$  {
        rollback;
    }
    else {
        read( $X$ );
         $R-TS(X) = \max(R-TS(X), TS(T_i))$ ;
    }
}
```



# Basic Time stamp protocol/alg

```
C if  $T_i$  requests write( $X$ ) {  
    if ( $TS(T_i) < R-TS(X) \parallel TS(T_i) < W-TS(X)$ ) {  
        rollback;  
    }  
    else {  
        write( $X$ );  
         $W-TS(X) = TS(T_i)$ ;  
    }  
}
```



# Strict Timestamp Ordering

---

## **1. Transaction T issues a write\_item(X) operation:**

If  $TS(T) > R-TS(X)$ , then delay T until the transaction T' that wrote or read X has terminated (committed or aborted).

## **2. Transaction T issues a read\_item(X) operation:**

If  $TS(T) > W-TS(X)$ , then delay T until the transaction T' that wrote or read X has terminated (committed or aborted).



# Thomas' Write Rule

---

**Observation:** if  $TS(T_i) < TS(T_j)$ , and  $T_i$  writes  $X$  before  $TS(T_j)$ , then we can just ignore that write, since it will be overwritten by  $T_j$  anyway

Thomas' write rule is a protocol tweak on the  $write(X)$  case based on the observation:

...

```
if  $TS(T_i) < W-TS(X)$  {  
  ignore write(X);  
}
```



# Resources

---

Chapter 22 of Fundamentals of Database Systems (FODS),  
6<sup>th</sup> Edition.

Internet Surf

(Most of the slides adapted from Ramez Elmasri and  
Shamkant Navathe , FODS)

